

titanic_survival

September 18, 2025

```
[1]: import os
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from collections import Counter

# sklearn imports
from sklearn.model_selection import train_test_split, cross_val_score, GridSearchCV
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report, roc_auc_score, roc_curve
import joblib

# plotting style
sns.set_theme(style="whitegrid")
%matplotlib inline
```

1 Load the data

```
[3]: TRAIN_PATH = "../data/train.csv"
TEST_PATH = "../data/test.csv"

# load
train = pd.read_csv(TRAIN_PATH)
test = pd.read_csv(TEST_PATH)

print("Train shape:", train.shape)
print("Test shape:", test.shape)
```

Train shape: (891, 12)

Test shape: (418, 11)

2 Quick look at the data (head, info, missing values)

```
[4]: display(train.head())
print("\n--- Info ---")
print(train.info())
print("\n--- Missing values (train) ---")
print(train.isnull().sum().sort_values(ascending=False).head(10))

print("\n--- Missing values (test) ---")
print(test.isnull().sum().sort_values(ascending=False).head(10))
```

	PassengerId	Survived	Pclass	\
0	1	0	3	
1	2	1	1	
2	3	1	3	
3	4	1	1	
4	5	0	3	

	Name	Sex	Age	SibSp	\
0	Braund, Mr. Owen Harris	male	22.0	1	
1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	
2	Heikkinen, Miss. Laina	female	26.0	0	
3	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	
4	Allen, Mr. William Henry	male	35.0	0	

	Parch	Ticket	Fare	Cabin	Embarked
0	0	A/5 21171	7.2500	NaN	S
1	0	PC 17599	71.2833	C85	C
2	0	STON/O2. 3101282	7.9250	NaN	S
3	0	113803	53.1000	C123	S
4	0	373450	8.0500	NaN	S

--- Info ---

```
<class 'pandas.core.frame.DataFrame'>
```

RangeIndex: 891 entries, 0 to 890

Data columns (total 12 columns):

#	Column	Non-Null Count	Dtype
0	PassengerId	891 non-null	int64
1	Survived	891 non-null	int64
2	Pclass	891 non-null	int64
3	Name	891 non-null	object
4	Sex	891 non-null	object
5	Age	714 non-null	float64
6	SibSp	891 non-null	int64

```

7   Parch      891 non-null    int64
8   Ticket     891 non-null    object
9   Fare       891 non-null    float64
10  Cabin      204 non-null    object
11  Embarked   889 non-null    object
dtypes: float64(2), int64(5), object(5)
memory usage: 83.7+ KB
None

```

--- Missing values (train) ---

```

Cabin      687
Age        177
Embarked    2
PassengerId 0
Name        0
Pclass      0
Survived    0
Sex         0
Parch       0
SibSp       0
dtype: int64

```

--- Missing values (test) ---

```

Cabin      327
Age         86
Fare        1
Name        0
Pclass      0
PassengerId 0
Sex         0
Parch       0
SibSp       0
Ticket      0
dtype: int64

```

3 4) Target distribution & basic EDA

Let's inspect survival balance and some univariate distributions.

```

[5]: # Target distribution
plt.figure(figsize=(6,4))
sns.countplot(x='Survived', data=train)
plt.title("Survived distribution (train)")
plt.xticks([0,1], ["Died (0)", "Survived (1)"])
plt.show()



# Age distribution

```

```
plt.figure(figsize=(8,4))
sns.histplot(train['Age'].dropna(), kde=True, bins=30)
plt.title("Age distribution (train)")
plt.show()

# Fare distribution (log-scale for visualization)
plt.figure(figsize=(8,4))
sns.histplot(train['Fare'].dropna(), bins=40)
plt.title("Fare distribution (train)")
plt.show()
```






4 5) Bivariate analysis: features vs target

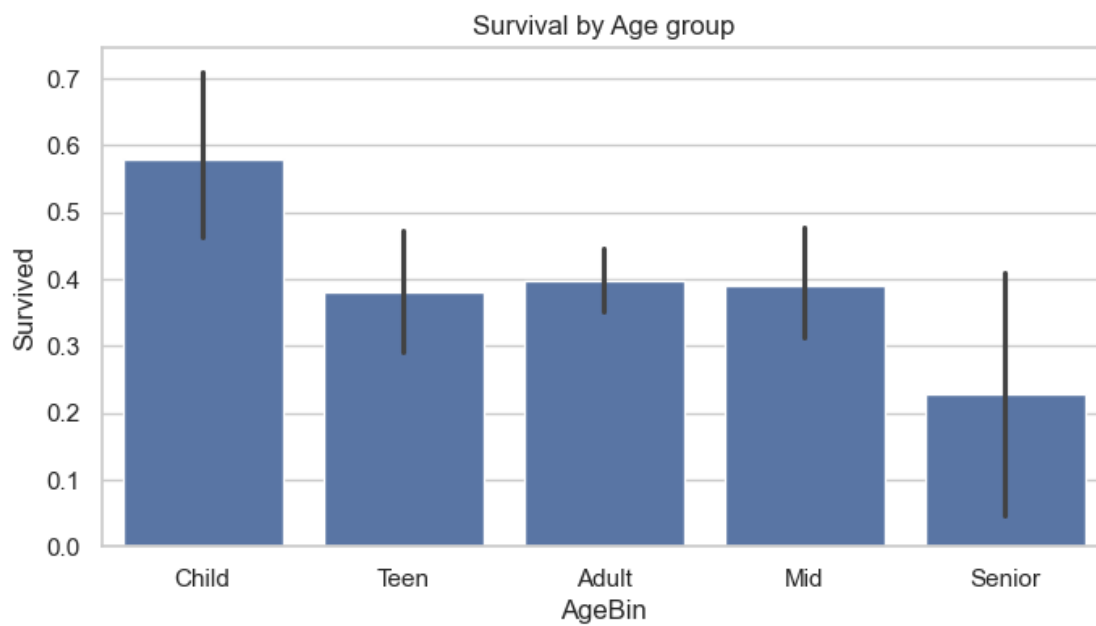

See how sex, class, embarcation relate to survival.

```
[6]: # Survival by Sex
plt.figure(figsize=(6,4))
sns.barplot(x='Sex', y='Survived', data=train)
plt.title("Survival rate by Sex")
plt.show()

# Survival by Pclass
plt.figure(figsize=(6,4))
sns.barplot(x='Pclass', y='Survived', data=train)
plt.title("Survival rate by Pclass")
plt.show()

# Survival vs Age (age bins)
train['AgeBin'] = pd.cut(train['Age'], bins=[0,12,20,40,60,80],
    ↳labels=['Child','Teen','Adult','Mid','Senior'])
plt.figure(figsize=(8,4))
sns.barplot(x='AgeBin', y='Survived', data=train,
    ↳order=['Child','Teen','Adult','Mid','Senior'])
plt.title("Survival by Age group")
plt.show()
train.drop(columns=['AgeBin'], inplace=True)
```





5 6) Feature engineering: create helpful features

We'll create: - Title from Name (Mr, Mrs, Miss, Master, Rare) - FamilySize = SibSp + Parch + 1 - IsAlone from FamilySize - Deck from Cabin (optional; many NA)

```
[7]: def extract_title(name):
    if pd.isna(name):
        return "None"
    # common format: "Last, Title. First"
    title = name.split(',')[1].split('.')[0].strip()
    return title

def simplify_title(title):
    # Map many rare titles to 'Rare' and standardize
    title = title.lower()
    if title in ['mr', 'mrs', 'miss', 'master']:
        return title.title()
    if title in ['ms']:
        return 'Miss'
    if title in ['mme', 'mademoiselle']:
        return 'Mrs'
    # everything else rare
    return 'Rare'

def add_features(df):
    df = df.copy()
    # Title
    df['Title'] = df['Name'].apply(extract_title).apply(simplify_title)
    # Family size
    df['FamilySize'] = df['SibSp'] + df['Parch'] + 1
    df['IsAlone'] = (df['FamilySize'] == 1).astype(int)
    # Deck (first letter of Cabin). Will be NaN if Cabin missing.
    df['Deck'] = df['Cabin'].apply(lambda x: str(x)[0] if pd.notna(x) else np.
    nan)
    return df

train_fe = add_features(train)
test_fe = add_features(test)

print("Titles in train:", train_fe['Title'].value_counts().to_dict())
display(train_fe[['Title', 'FamilySize', 'IsAlone', 'Deck']].head())
```

Titles in train: {'Mr': 517, 'Miss': 183, 'Mrs': 126, 'Master': 40, 'Rare': 25}

	Title	FamilySize	IsAlone	Deck
0	Mr	2	0	NaN
1	Mrs	2	0	C
2	Miss	1	1	NaN
3	Mrs	2	0	C

4 Mr 1 1 NaN

6 7) Handle missing values with domain-specific logic

Strategy: - Fill Embarked with mode - Fill Fare missing in test with median - Impute Age by median of Title groups when available, else overall median - Drop Cabin (too many missing) — we keep Deck if desired but it has lots of missing

```
[9]: # Copy dataframes to avoid modifying originals
train_clean = train_fe.copy()
test_clean = test_fe.copy()

# Embarked -> fill with mode
for df in [train_clean, test_clean]:
    if 'Embarked' in df.columns:
        df['Embarked'] = df['Embarked'].fillna(df['Embarked'].mode()[0])

# Fare -> fill with median (test may have missing Fare)
test_clean['Fare'] = test_clean['Fare'].fillna(test_clean['Fare'].median())

# Age -> fill by Title median
title_age_median = train_clean.groupby('Title')['Age'].median()
# fallback median
overall_age_median = train_clean['Age'].median()

def fill_age_by_title(row):
    if pd.isna(row['Age']):
        return row['Age']
    title = row['Title']
    if pd.isna(title) and title in title_age_median.index and pd.
    isna(title_age_median.loc[title]):
        return title_age_median.loc[title]
    return overall_age_median

train_clean['Age'] = train_clean.apply(fill_age_by_title, axis=1)
test_clean['Age'] = test_clean.apply(fill_age_by_title, axis=1)

# Drop Cabin column (we have Deck but it's sparse)
for df in [train_clean, test_clean]:
    if 'Cabin' in df.columns:
        df.drop(columns=['Cabin'], inplace=True)

# Quick check
print("Missing values after cleaning (train):")
print(train_clean.isnull().sum().sort_values(ascending=False).head(10))
print("\nMissing values after cleaning (test):")
print(test_clean.isnull().sum().sort_values(ascending=False).head(10))
```

Missing values after cleaning (train):

```
Deck          687
Survived       0
PassengerId    0
Name           0
Sex            0
Age            0
Pclass         0
SibSp          0
Parch          0
Fare           0
dtype: int64
```

Missing values after cleaning (test):

```
Deck          327
PassengerId    0
Name           0
Sex            0
Age            0
Pclass         0
SibSp          0
Parch          0
Fare           0
Ticket         0
dtype: int64
```

7 8) Prepare final feature list & split target

We'll select a set of features (mix of numeric & categorical) and prepare X/y.

```
[10]: # Feature selection
features = ['Pclass', 'Sex', 'Age', 'SibSp', 'Parch', 'Fare', 'Embarked',
           ↪ 'Title', 'FamilySize', 'IsAlone', 'Deck']
# Note: Deck will have many NaNs; it will be handled by the imputer/encoder

X = train_clean[features].copy()
y = train_clean['Survived'].copy()

X_test_final = test_clean[features].copy() # for final predictions

print("X shape:", X.shape)
display(X.head())
```

X shape: (891, 11)

	Pclass	Sex	Age	SibSp	Parch	Fare	Embarked	Title	FamilySize	\
0	3	male	22.0	1	0	7.2500	S	Mr	2	
1	1	female	38.0	1	0	71.2833	C	Mrs	2	

2	3	female	26.0	0	0	7.9250	S	Miss	1
3	1	female	35.0	1	0	53.1000	S	Mrs	2
4	3	male	35.0	0	0	8.0500	S	Mr	1

	IsAlone	Deck
0	0	NaN
1	0	C
2	1	NaN
3	0	C
4	1	NaN

8 9) Split into training and validation sets

We'll use an 80/20 split for local evaluation (stratify by y).

```
[11]: X_train, X_val, y_train, y_val = train_test_split(
      X, y, test_size=0.20, random_state=42, stratify=y
    )
    print("Train:", X_train.shape, "Validation:", X_val.shape)
```

Train: (712, 11) Validation: (179, 11)

9 10) Build preprocessing pipeline

We will: - Impute numeric columns by median and scale them - Impute categorical columns by most frequent and one-hot encode them - Combine using `ColumnTransformer` so the pipeline can be used on train & test equally

```
[13]: from sklearn.pipeline import Pipeline
      from sklearn.preprocessing import OneHotEncoder
      from sklearn.impute import SimpleImputer
      from sklearn.compose import ColumnTransformer

      numeric_features = ['Age', 'SibSp', 'Parch', 'Fare', 'FamilySize']
      categorical_features = ['Pclass', 'Sex', 'Embarked', 'Title', 'IsAlone', 'Deck']

      numeric_transformer = Pipeline(steps=[
          ('imputer', SimpleImputer(strategy='median')),
          ('scaler', StandardScaler())
      ])

      categorical_transformer = Pipeline(steps=[
          ('imputer', SimpleImputer(strategy='most_frequent')),
          ('onehot', OneHotEncoder(handle_unknown='ignore', sparse_output=False))
      ])

      preprocessor = ColumnTransformer(transformers=[
```

```

        ('num', numeric_transformer, numeric_features),
        ('cat', categorical_transformer, categorical_features)
    ], remainder='drop')

# Test that preprocessor can fit-transform training data
preprocessor.fit(X_train)
print("Preprocessor fitted.")

```

Preprocessor fitted.

10 11) Baseline model pipelines

We'll create two pipelines: - Logistic Regression (simple, interpretable) - Random Forest (powerful tree-based)

```

[15]: # Logistic Regression pipeline
log_pipe = Pipeline(steps=[
    ('preprocessor', preprocessor),
    ('clf', LogisticRegression(max_iter=1000, random_state=42))
])

# Random Forest pipeline
rf_pipe = Pipeline(steps=[
    ('preprocessor', preprocessor),
    ('clf', RandomForestClassifier(n_estimators=100, random_state=42))
])

# Quick cross-validation (5-fold) to get baseline scores
from sklearn.model_selection import cross_val_score
print("Logistic CV:", np.mean(cross_val_score(log_pipe, X_train, y_train, cv=5,
    ↪scoring='accuracy')))
print("RandomForest CV:", np.mean(cross_val_score(rf_pipe, X_train, y_train,
    ↪cv=5, scoring='accuracy')))

```

Logistic CV: 0.8188811188811188

RandomForest CV: 0.7922387471683245

11 12) Train on training set and evaluate on validation set

Fit both pipelines and compute accuracy, classification report, confusion matrix, and ROC AUC.

```

[16]: # Fit Logistic Regression
log_pipe.fit(X_train, y_train)
y_pred_log = log_pipe.predict(X_val)
y_prob_log = log_pipe.predict_proba(X_val)[: ,1]

print("Logistic Regression")

```

```

print("Accuracy:", accuracy_score(y_val, y_pred_log))
print(classification_report(y_val, y_pred_log))
print("ROC AUC:", roc_auc_score(y_val, y_prob_log))
print("Confusion Matrix:\n", confusion_matrix(y_val, y_pred_log))

# Fit Random Forest
rf_pipe.fit(X_train, y_train)
y_pred_rf = rf_pipe.predict(X_val)
y_prob_rf = rf_pipe.predict_proba(X_val)[: ,1]

print("\nRandom Forest")
print("Accuracy:", accuracy_score(y_val, y_pred_rf))
print(classification_report(y_val, y_pred_rf))
print("ROC AUC:", roc_auc_score(y_val, y_prob_rf))
print("Confusion Matrix:\n", confusion_matrix(y_val, y_pred_rf))

```

Logistic Regression

Accuracy: 0.8603351955307262

	precision	recall	f1-score	support
0	0.87	0.91	0.89	110
1	0.84	0.78	0.81	69
accuracy			0.86	179
macro avg	0.86	0.85	0.85	179
weighted avg	0.86	0.86	0.86	179

ROC AUC: 0.8735177865612647

Confusion Matrix:

```

[[100 10]
 [ 15 54]]

```

Random Forest

Accuracy: 0.7988826815642458

	precision	recall	f1-score	support
0	0.82	0.86	0.84	110
1	0.76	0.70	0.73	69
accuracy			0.80	179
macro avg	0.79	0.78	0.78	179
weighted avg	0.80	0.80	0.80	179

ROC AUC: 0.8228590250329382

Confusion Matrix:

```

[[95 15]
 [21 48]]


```

12 13) ROC curve comparison

Plot ROC curves for both models on the validation set.

```
[17]: fpr_log, tpr_log, _ = roc_curve(y_val, y_prob_log)
      fpr_rf, tpr_rf, _ = roc_curve(y_val, y_prob_rf)

      plt.figure(figsize=(8,6))
      plt.plot(fpr_log, tpr_log, label=f'Logistic_
      ↪(AUC={roc_auc_score(y_val,y_prob_log):.3f})')
      plt.plot(fpr_rf, tpr_rf, label=f'RandomForest_
      ↪(AUC={roc_auc_score(y_val,y_prob_rf):.3f})')
      plt.plot([0,1],[0,1], 'k--')
      plt.xlabel("False Positive Rate")
      plt.ylabel("True Positive Rate")
      plt.title("ROC Curves")
      plt.legend()
      plt.show()
```



13 14) Feature importance (Random Forest)

To show feature importances we need the post-preprocessor feature names and the classifier's `feature_importances_`.

```
[18]: # Extract numeric feature names (they remain as is)
num_feats = numeric_features

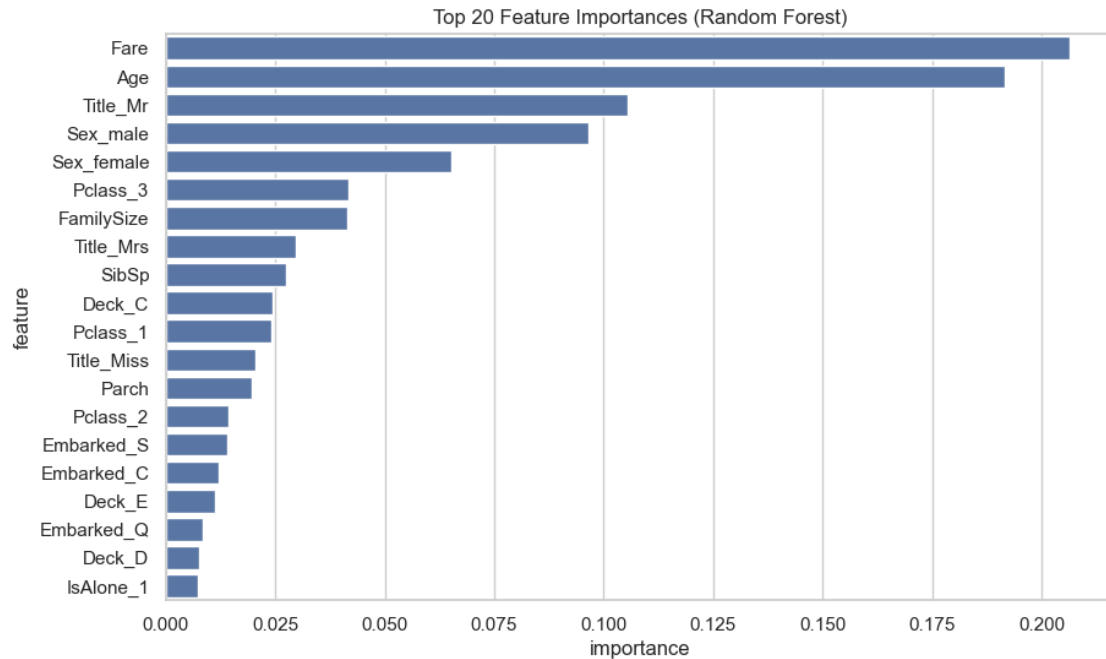
# Extract categorical feature names after OneHot encoding
# We need to get fitted OneHotEncoder to obtain categories
ohe = rf_pipe.named_steps['preprocessor'].named_transformers_['cat'].
    ↪named_steps['onehot']
try:
    cat_feature_names = list(ohe.get_feature_names_out(categorical_features))
except:
    # fallback for older sklearn versions
    cat_feature_names = []
    for i, cat in enumerate(categorical_features):
        cats = ohe.categories_[i]
        cat_feature_names += [f"{cat}_{c}" for c in cats]

feature_names = num_feats + cat_feature_names

# Get importances
importances = rf_pipe.named_steps['clf'].feature_importances_

# Build DataFrame and plot
feat_imp_df = pd.DataFrame({'feature': feature_names, 'importance':
    ↪importances})
feat_imp_df = feat_imp_df.sort_values(by='importance', ascending=False).
    ↪reset_index(drop=True)
plt.figure(figsize=(10,6))
sns.barplot(x='importance', y='feature', data=feat_imp_df.head(20))
plt.title("Top 20 Feature Importances (Random Forest)")
plt.show()

display(feat_imp_df.head(30))
```



	feature	importance
0	Fare	0.206273
1	Age	0.191500
2	Title_Mr	0.105510
3	Sex_male	0.096514
4	Sex_female	0.065112
5	Pclass_3	0.041688
6	FamilySize	0.041341
7	Title_Mrs	0.029868
8	SibSp	0.027396
9	Deck_C	0.024355
10	Pclass_1	0.024215
11	Title_Miss	0.020396
12	Parch	0.019639
13	Pclass_2	0.014462
14	Embarked_S	0.014108
15	Embarked_C	0.012118
16	Deck_E	0.011274
17	Embarked_Q	0.008531
18	Deck_D	0.007618
19	IsAlone_1	0.007444
20	Title_Master	0.006981
21	IsAlone_0	0.006648
22	Title_Rare	0.004793
23	Deck_B	0.004516
24	Deck_A	0.003098

25	Deck_F	0.002203
26	Deck_G	0.002032
27	Deck_T	0.000368

14 15) Hyperparameter tuning (Random Forest) — small Grid-SearchCV

This is a small grid; you can expand later if you want. We'll tune `n_estimators`, `max_depth`, and `min_samples_split`.

```
[19]: param_grid = {
        'clf__n_estimators': [100, 200],
        'clf__max_depth': [None, 6, 10],
        'clf__min_samples_split': [2, 5]
    }

    # Use GridSearchCV with a pipeline; scoring by accuracy
    grid_search = GridSearchCV(rf_pipe, param_grid, cv=5, scoring='accuracy',
                               n_jobs=-1, verbose=1)
    grid_search.fit(X_train, y_train)

    print("Best params:", grid_search.best_params_)
    print("Best CV score:", grid_search.best_score_)

    best_rf = grid_search.best_estimator_
    # evaluate on validation set
    y_pred_best = best_rf.predict(X_val)
    y_prob_best = best_rf.predict_proba(X_val)[:,1]
    print("Validation accuracy (best RF):", accuracy_score(y_val, y_pred_best))
    print("Validation ROC AUC (best RF):", roc_auc_score(y_val, y_prob_best))
```

```
Fitting 5 folds for each of 12 candidates, totalling 60 fits
Best params: {'clf__max_depth': 6, 'clf__min_samples_split': 5,
'clf__n_estimators': 200}
Best CV score: 0.8245149216980202
Validation accuracy (best RF): 0.8100558659217877
Validation ROC AUC (best RF): 0.8538208168642951
```

15 16) Train final model on the *entire* training dataset

Retrain the best model on full train data (not the local validation split) for final predictions on `test.csv`.

```
[20]: # Prepare full training data (we already have X, y from earlier)
    # If we used grid_search.best_estimator_, we can fit it on full data
    final_model = grid_search.best_estimator_
```

```
# Fit on full training set
final_model.fit(X, y)
print("Final model trained on full training data.")
```

Final model trained on full training data.

16 17) Create predictions for test.csv and build submission file

Kaggle expects a CSV with PassengerId and Survived columns.

```
[25]: # Ensure test has PassengerId
if 'PassengerId' not in test.columns:
    raise ValueError("test.csv must contain 'PassengerId' column for creating_
    ↪submission file.")

# Predict
test_preds = final_model.predict(X_test_final)

submission = pd.DataFrame({
    "PassengerId": test['PassengerId'],
    "Survived": test_preds.astype(int)
})

submission_filename = "../data/titanic_submission.csv"
submission.to_csv(submission_filename, index=False)
print(f"Submission saved to {submission_filename}")
display(submission.head())
```

Submission saved to ../data/titanic_submission.csv

	PassengerId	Survived
0	892	0
1	893	0
2	894	0
3	895	0
4	896	1

17 18) Save the final model to disk

We save the whole pipeline so that preprocessing is included.

```
[24]: model_filename = "../model/titanic_final_model.joblib"
joblib.dump(final_model, model_filename)
print(f"Saved model pipeline to {model_filename}")
```

Saved model pipeline to ../model/titanic_final_model.joblib