

titanic_survival

October 15, 2025

1 1) Importing Libraries

```
[78]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from collections import Counter

# sklearn imports
from sklearn.model_selection import train_test_split, cross_val_score, GridSearchCV
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report, roc_auc_score, roc_curve
import joblib

# plotting style
sns.set_theme(style="whitegrid")
%matplotlib inline
```

2 2) Load the data

```
[79]: TRAIN_PATH = "../data/train.csv"
TEST_PATH = "../data/test.csv"

train = pd.read_csv(TRAIN_PATH)
test = pd.read_csv(TEST_PATH)

print("Train shape:", train.shape)
print("Test shape:", test.shape)
```

Train shape: (891, 12)

Test shape: (418, 11)

3) Quick look at the data (head, info, missing values)

```
[80]: display(train.head())
print("\nInfo")
print(train.info())
print("\nMissing values (train)")
print(train.isnull().sum().sort_values(ascending=False).head(10))

print("\nMissing values (test)")
print(test.isnull().sum().sort_values(ascending=False).head(10))
```

	PassengerId	Survived	Pclass	\
0	1	0	3	
1	2	1	1	
2	3	1	3	
3	4	1	1	
4	5	0	3	

	Name	Sex	Age	SibSp	\
0	Braund, Mr. Owen Harris	male	22.0	1	
1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	
2	Heikkinen, Miss. Laina	female	26.0	0	
3	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	
4	Allen, Mr. William Henry	male	35.0	0	

	Parch	Ticket	Fare	Cabin	Embarked
0	0	A/5 21171	7.2500	NaN	S
1	0	PC 17599	71.2833	C85	C
2	0	STON/O2. 3101282	7.9250	NaN	S
3	0	113803	53.1000	C123	S
4	0	373450	8.0500	NaN	S

Info

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 891 entries, 0 to 890
```

```
Data columns (total 12 columns):
```

#	Column	Non-Null Count	Dtype
0	PassengerId	891 non-null	int64
1	Survived	891 non-null	int64
2	Pclass	891 non-null	int64
3	Name	891 non-null	object
4	Sex	891 non-null	object
5	Age	714 non-null	float64
6	SibSp	891 non-null	int64
7	Parch	891 non-null	int64
8	Ticket	891 non-null	object

```

9   Fare            891 non-null   float64
10  Cabin           204 non-null   object
11  Embarked       889 non-null   object
dtypes: float64(2), int64(5), object(5)
memory usage: 83.7+ KB
None

```

Missing values (train)

```

Cabin           687
Age             177
Embarked        2
PassengerId     0
Name            0
Pclass          0
Survived        0
Sex             0
Parch           0
SibSp           0
dtype: int64

```

Missing values (test)

```

Cabin           327
Age             86
Fare            1
Name            0
Pclass          0
PassengerId     0
Sex             0
Parch           0
SibSp           0
Ticket          0
dtype: int64

```

4 4) Target distribution & basic EDA

Inspect survival balance and some univariate distributions.

```

[81]: # Target distribution
plt.figure(figsize=(6,4))
sns.countplot(x='Survived', data=train)
plt.title("Survived distribution (train)")
plt.xticks([0,1], ["Died (0)", "Survived (1)"])
plt.show()



# Age distribution
plt.figure(figsize=(8,4))
sns.histplot(train['Age'].dropna(), kde=True, bins=30)

```

```
plt.title("Age distribution (train)")
plt.show()

# Fare distribution (log-scale for visualization)
plt.figure(figsize=(8,4))
sns.histplot(train['Fare'].dropna(), bins=40)
plt.title("Fare distribution (train)")
plt.show()
```






5 5) Bivariate analysis: features vs target


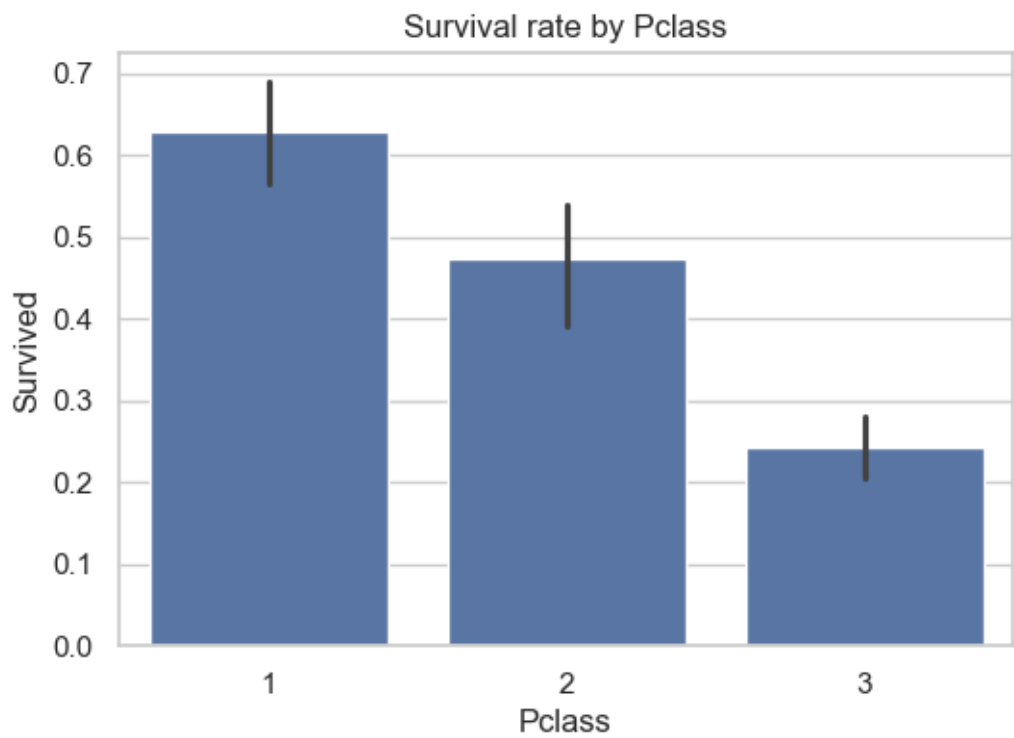
See how sex, class, embarcation relate to survival.

```
[82]: # Survival by Sex
plt.figure(figsize=(6,4))
sns.barplot(x='Sex', y='Survived', data=train)
plt.title("Survival rate by Sex")
plt.show()

# Survival by Pclass
plt.figure(figsize=(6,4))
sns.barplot(x='Pclass', y='Survived', data=train)
plt.title("Survival rate by Pclass")
plt.show()

# Survival vs Age (age bins)
train['AgeBin'] = pd.cut(train['Age'], bins=[0,12,20,40,60,80],
    ↳labels=['Child','Teen','Adult','Mid','Senior'])
plt.figure(figsize=(8,4))
sns.barplot(x='AgeBin', y='Survived', data=train,
    ↳order=['Child','Teen','Adult','Mid','Senior'])
plt.title("Survival by Age group")
plt.show()
train.drop(columns=['AgeBin'], inplace=True)
```





6 6) Feature engineering: create helpful features

- Title from Name (Mr, Mrs, Miss, Master, Rare)
- FamilySize = SibSp + Parch + 1
- IsAlone from FamilySize

```
[83]: def extract_title(name):
    if pd.isna(name):
        return "None"
    title = name.split(',')[1].split('.')[0].strip()
    return title

def simplify_title(title):
    title = title.lower()
    if title in ['mr', 'mrs', 'miss', 'master']:
        return title.title()
    if title in ['ms']:
        return 'Miss'
    if title in ['mme', 'mademoiselle']:
        return 'Mrs'
    return 'Rare'

def add_features(df):
    df = df.copy()
    df['Title'] = df['Name'].apply(extract_title).apply(simplify_title)
    df['FamilySize'] = df['SibSp'] + df['Parch'] + 1
    df['IsAlone'] = (df['FamilySize'] == 1).astype(int)
    return df

train_fe = add_features(train)
test_fe = add_features(test)

print("Titles in train:", train_fe['Title'].value_counts())
display(train_fe[['Title', 'FamilySize', 'IsAlone']].head())
```

Titles in train: Title

Mr	517
Miss	183
Mrs	126
Master	40
Rare	25

Name: count, dtype: int64

	Title	FamilySize	IsAlone
0	Mr	2	0
1	Mrs	2	0
2	Miss	1	1
3	Mrs	2	0
4	Mr	1	1

7 7) Handle missing values with domain-specific logic

```
[84]: train_clean = train_fe.copy()
test_clean = test_fe.copy()

# Embarked -> fill with mode
for df in [train_clean, test_clean]:
    if 'Embarked' in df.columns:
        df['Embarked'] = df['Embarked'].fillna(df['Embarked'].mode()[0])

# Fare -> fill with median (test may have missing Fare)
test_clean['Fare'] = test_clean['Fare'].fillna(test_clean['Fare'].median())

# Age -> fill by Title median
title_age_median = train_clean.groupby('Title')['Age'].median()
# fallback median
overall_age_median = train_clean['Age'].median()

def fill_age_by_title(row):
    if pd.isna(row['Age']):
        return row['Age']
    title = row['Title']
    if pd.isna(title) and title in title_age_median.index and pd.
    isna(title_age_median.loc[title]):
        return title_age_median.loc[title]
    return overall_age_median

train_clean['Age'] = train_clean.apply(fill_age_by_title, axis=1)
test_clean['Age'] = test_clean.apply(fill_age_by_title, axis=1)

for df in [train_clean, test_clean]:
    if 'Cabin' in df.columns:
        df.drop(columns=['Cabin'], inplace=True)

# Quick check
print("Missing values after cleaning (train):")
print(train_clean.isnull().sum().sort_values(ascending=False).head(10))
print("\nMissing values after cleaning (test):")
print(test_clean.isnull().sum().sort_values(ascending=False).head(10))
```

Missing values after cleaning (train):

PassengerId	0
Survived	0
Pclass	0
Name	0
Sex	0
Age	0
SibSp	0

```
Parch      0
Ticket     0
Fare       0
dtype: int64
```

Missing values after cleaning (test):

```
PassengerId  0
Pclass       0
Name         0
Sex          0
Age         0
SibSp       0
Parch       0
Ticket      0
Fare        0
Embarked    0
dtype: int64
```

8 8) Prepare final feature list & split target

We'll select a set of features (mix of numeric & categorical) and prepare X/y.

```
[85]: # Feature selection
features = ['Pclass', 'Sex', 'Age', 'SibSp', 'Parch', 'Fare', 'Embarked',
           ↪ 'Title', 'FamilySize', 'IsAlone']

X = train_clean[features].copy()
y = train_clean['Survived'].copy()

X_test_final = test_clean[features].copy() # for final predictions

print("X shape:", X.shape)
display(X.head())
```

X shape: (891, 10)

	Pclass	Sex	Age	SibSp	Parch	Fare	Embarked	Title	FamilySize	\
0	3	male	22.0	1	0	7.2500	S	Mr	2	
1	1	female	38.0	1	0	71.2833	C	Mrs	2	
2	3	female	26.0	0	0	7.9250	S	Miss	1	
3	1	female	35.0	1	0	53.1000	S	Mrs	2	
4	3	male	35.0	0	0	8.0500	S	Mr	1	

	IsAlone
0	0
1	0
2	1
3	0

9) Split into training and validation sets

We'll use an 80/20 split for local evaluation (stratify by y).

```
[86]: X_train, X_val, y_train, y_val = train_test_split(
        X, y, test_size=0.20, random_state=42, stratify=y
    )
    print("Train:", X_train.shape, "Validation:", X_val.shape)
```

Train: (712, 10) Validation: (179, 10)

10) Build preprocessing pipeline

We will: - Impute numeric columns by median and scale them - Impute categorical columns by most frequent and one-hot encode them - Combine using `ColumnTransformer` so the pipeline can be used on train & test equally

```
[87]: from sklearn.pipeline import Pipeline
    from sklearn.preprocessing import OneHotEncoder
    from sklearn.impute import SimpleImputer
    from sklearn.compose import ColumnTransformer

    numeric_features = ['Age', 'SibSp', 'Parch', 'Fare', 'FamilySize']
    categorical_features = ['Pclass', 'Sex', 'Embarked', 'Title', 'IsAlone']

    numeric_transformer = Pipeline(steps=[
        ('imputer', SimpleImputer(strategy='median')),
        ('scaler', StandardScaler())
    ])

    categorical_transformer = Pipeline(steps=[
        ('imputer', SimpleImputer(strategy='most_frequent')),
        ('onehot', OneHotEncoder(handle_unknown='ignore', sparse_output=False))
    ])

    preprocessor = ColumnTransformer(transformers=[
        ('num', numeric_transformer, numeric_features),
        ('cat', categorical_transformer, categorical_features)
    ], remainder='drop')

    preprocessor.fit(X_train)
```

```
[87]: ColumnTransformer(transformers=[('num',
                                         Pipeline(steps=[('imputer',
                                                             SimpleImputer(strategy='median'))],
```

```

        ('scaler', StandardScaler()))],
        ['Age', 'SibSp', 'Parch', 'Fare',
         'FamilySize']),
        ('cat',
         Pipeline(steps=[('imputer',
                           SimpleImputer(strategy='most_frequent')),
                           ('onehot',
                            OneHotEncoder(handle_unknown='ignore',
                                           sparse_output=False))])),
        ['Pclass', 'Sex', 'Embarked', 'Title',
         'IsAlone']]))

```

11 11) Baseline model pipelines

We'll create two pipelines: - Logistic Regression (simple, interpretable) - Random Forest (powerful tree-based)

```

[88]: # Logistic Regression pipeline
log_pipe = Pipeline(steps=[
    ('preprocessor', preprocessor),
    ('clf', LogisticRegression(max_iter=1000, random_state=42))
])

# Random Forest pipeline
rf_pipe = Pipeline(steps=[
    ('preprocessor', preprocessor),
    ('clf', RandomForestClassifier(n_estimators=100, random_state=42))
])

```

Quick cross-validation (5-fold) to get baseline scores

```

[89]: from sklearn.model_selection import cross_val_score
print("Logistic CV:", np.mean(cross_val_score(log_pipe, X_train, y_train, cv=5,
↪scoring='accuracy'))))
print("RandomForest CV:", np.mean(cross_val_score(rf_pipe, X_train, y_train,
↪cv=5, scoring='accuracy'))))

```

```

Logistic CV: 0.8202698709740963
RandomForest CV: 0.7894316950654978

```

12 12) Train on training set and evaluate on validation set

Fit both pipelines and compute accuracy, classification report, confusion matrix, and ROC AUC.

```

[90]: # Fit Logistic Regression
log_pipe.fit(X_train, y_train)
y_pred_log = log_pipe.predict(X_val)

```

```

y_prob_log = log_pipe.predict_proba(X_val)[: ,1]

print("Logistic Regression")
print("Accuracy:", accuracy_score(y_val, y_pred_log))
print(classification_report(y_val, y_pred_log))
print("ROC AUC:", roc_auc_score(y_val, y_prob_log))
print("Confusion Matrix:\n", confusion_matrix(y_val, y_pred_log))

# Fit Random Forest
rf_pipe.fit(X_train, y_train)
y_pred_rf = rf_pipe.predict(X_val)
y_prob_rf = rf_pipe.predict_proba(X_val)[: ,1]

print("\nRandom Forest")
print("Accuracy:", accuracy_score(y_val, y_pred_rf))
print(classification_report(y_val, y_pred_rf))
print("ROC AUC:", roc_auc_score(y_val, y_prob_rf))
print("Confusion Matrix:\n", confusion_matrix(y_val, y_pred_rf))

```

Logistic Regression

Accuracy: 0.8547486033519553

	precision	recall	f1-score	support
0	0.86	0.91	0.88	110
1	0.84	0.77	0.80	69
accuracy			0.85	179
macro avg	0.85	0.84	0.84	179
weighted avg	0.85	0.85	0.85	179

ROC AUC: 0.8782608695652174

Confusion Matrix:

```

[[100  10]
 [ 16  53]]

```

Random Forest

Accuracy: 0.8156424581005587

	precision	recall	f1-score	support
0	0.83	0.87	0.85	110
1	0.78	0.72	0.75	69
accuracy			0.82	179
macro avg	0.81	0.80	0.80	179
weighted avg	0.81	0.82	0.81	179

ROC AUC: 0.8276021080368907

Confusion Matrix:


```
[[96 14]
 [19 50]]
```

13 13) ROC curve comparison

Plot ROC curves for both models on the validation set.

```
[91]: fpr_log, tpr_log, _ = roc_curve(y_val, y_prob_log)
      fpr_rf, tpr_rf, _ = roc_curve(y_val, y_prob_rf)

      plt.figure(figsize=(8,6))
      plt.plot(fpr_log, tpr_log, label=f'Logistic_
      ↳(AUC={roc_auc_score(y_val,y_prob_log):.3f})')
      plt.plot(fpr_rf, tpr_rf, label=f'RandomForest_
      ↳(AUC={roc_auc_score(y_val,y_prob_rf):.3f})')
      plt.plot([0,1],[0,1], 'k--')
      plt.xlabel("False Positive Rate")
      plt.ylabel("True Positive Rate")
      plt.title("ROC Curves")
      plt.legend()
      plt.show()
```



14 14) Feature importance (Random Forest)

To show feature importances we need the post-preprocessor feature names and the classifier's `feature_importances_`.


```
[92]: num_feats = numeric_features
      ohe = rf_pipe.named_steps['preprocessor'].named_transformers_['cat'].
          ↪named_steps['onehot']

      try:
          cat_feature_names = list(ohe.get_feature_names_out(categorical_features))
      except:
          cat_feature_names = []
          for i, cat in enumerate(categorical_features):
              cats = ohe.categories_[i]
              cat_feature_names += [f"{cat}_{c}" for c in cats]

      feature_names = num_feats + cat_feature_names

      # Get importances
      importances = rf_pipe.named_steps['clf'].feature_importances_

      # Build DataFrame and plot
      feat_imp_df = pd.DataFrame({'feature': feature_names, 'importance': ↪
          ↪importances})
      feat_imp_df = feat_imp_df.sort_values(by='importance', ascending=False).
          ↪reset_index(drop=True)
      plt.figure(figsize=(10,6))
      sns.barplot(x='importance', y='feature', data=feat_imp_df.head(20))
      plt.title("Top 20 Feature Importances (Random Forest)")
      plt.show()
```



15 15) Hyperparameter tuning (Random Forest) — small Grid-SearchCV

```
[93]: param_grid = {
    'clf__n_estimators': [100, 200],
    'clf__max_depth': [None, 6, 10],
    'clf__min_samples_split': [2, 5]
}

# Use GridSearchCV with a pipeline; scoring by accuracy
grid_search = GridSearchCV(rf_pipe, param_grid, cv=5, scoring='accuracy',
    ↪ n_jobs=-1, verbose=1)
grid_search.fit(X_train, y_train)

print("Best params:", grid_search.best_params_)
print("Best CV score:", grid_search.best_score_)

best_rf = grid_search.best_estimator_
# evaluate on validation set
y_pred_best = best_rf.predict(X_val)
y_prob_best = best_rf.predict_proba(X_val)[: ,1]
print("Validation accuracy (best RF):", accuracy_score(y_val, y_pred_best))
print("Validation ROC AUC (best RF):", roc_auc_score(y_val, y_prob_best))
```

Fitting 5 folds for each of 12 candidates, totalling 60 fits


```
Best params: {'clf__max_depth': 10, 'clf__min_samples_split': 5,
'clf__n_estimators': 100}
Best CV score: 0.8217078695951937
Validation accuracy (best RF): 0.8324022346368715
Validation ROC AUC (best RF): 0.8432806324110672
```

```
[94]: grid_search.best_estimator_.fit(X_train, y_train)
y_pred_rf = grid_search.best_estimator_.predict(X_val)
y_prob_rf = grid_search.best_estimator_.predict_proba(X_val)[:,:1]

print("\nRandom Forest")
print("Accuracy:", accuracy_score(y_val, y_pred_rf))
print(classification_report(y_val, y_pred_rf))
print("ROC AUC:", roc_auc_score(y_val, y_prob_rf))
print("Confusion Matrix:\n", confusion_matrix(y_val, y_pred_rf))
```

Random Forest

Accuracy: 0.8324022346368715

	precision	recall	f1-score	support
0	0.83	0.91	0.87	110
1	0.83	0.71	0.77	69
accuracy			0.83	179
macro avg	0.83	0.81	0.82	179
weighted avg	0.83	0.83	0.83	179

ROC AUC: 0.8432806324110672

Confusion Matrix:

```
[[100  10]
 [ 20  49]]
```

16 16) Train both modes on the entire training dataset

Retrain the best model on full train data (not the local validation split) for final predictions on test.csv.

```
[95]: log_model = log_pipe
rf_model=grid_search.best_estimator_

# Fit on full training set
log_model.fit(X, y)
rf_model.fit(X, y)

print("Models trained on full training data.")
```

Models trained on full training data.

17 17) Create predictions for test.csv and build submission file

Kaggle expects a CSV with PassengerId and Survived columns.

```
[96]: # Ensure test has PassengerId
if 'PassengerId' not in test.columns:
    raise ValueError("test.csv must contain 'PassengerId' column for creating_
submission files.")

#Logistic Regression Submission
log_preds = log_model.predict(X_test_final)

submission_log = pd.DataFrame({
    "PassengerId": test['PassengerId'],
    "Survived": log_preds.astype(int)
})

submission_log_filename = "../data/submission_log.csv"
submission_log.to_csv(submission_log_filename, index=False)
print(f"Logistic Regression submission saved to {submission_log_filename}")
display(submission_log.head())

#Random Forest Submission
rf_preds = rf_model.predict(X_test_final)

submission_rf = pd.DataFrame({
    "PassengerId": test['PassengerId'],
    "Survived": rf_preds.astype(int)
})

submission_rf_filename = "../data/submission_rf.csv"
submission_rf.to_csv(submission_rf_filename, index=False)
print(f"Random Forest submission saved to {submission_rf_filename}")
display(submission_rf.head())
```

Logistic Regression submission saved to ../data/submission_log.csv

	PassengerId	Survived
0	892	0
1	893	1
2	894	0
3	895	0
4	896	1

Random Forest submission saved to ../data/submission_rf.csv

	PassengerId	Survived
--	-------------	----------

0	892	0
1	893	0
2	894	0
3	895	0
4	896	1

18 18) Save the final model to disk

Save the whole pipeline so that preprocessing is included.

```
[97]: model_filename_log = "../model/titanic_log_model.joblib"
      model_filename_rf = "../model/titanic_rf_model.joblib"
      joblib.dump(log_model, model_filename_log)
      joblib.dump(rf_model, model_filename_rf)
      print(f"Saved models pipeline to {model_filename_log}")
      print(f"Saved models pipeline to {model_filename_rf}")
```

Saved models pipeline to ../model/titanic_log_model.joblib

Saved models pipeline to ../model/titanic_rf_model.joblib